

**Deccan Education Society's
Kirti M. Doongursee College of Arts, Science and Commerce
[Autonomous], Dadar (West), Mumbai-400 028.**



M.Sc. [Computer Science]

Practical journal

Seat Number []

Department of Computer Science and Information Technology

**Department of Computer Science and Information Technology
Deccan Education Society's**

**Kirti M. Doongursee College of Arts, Science and Commerce
[Autonomous], Dadar (West), Mumbai-400 028.**

C E R T I F I C A T E

This is to certify that _____

of M.Sc. (Computer Science) with Seat No. _____ has completed **Practical journal**
of Paper- **Applied Machine Learning and Deep Learning** under my supervision
in this College during the year **2023-2024**.

Lecturer-In-Charge

Date: / /2024

Examined by:

Date:

H.O.D.

**Department of
Computer Science & IT**

Date:

Remarks:

INDEX

| Sr. No. | Date | Title | Sign |
|---------|----------|---|------|
| 1. | 14/05/24 | Practical 1.1 linear regression on Study dataset | |
| 2. | 14/05/24 | Practical 1.2 linear regression on Advertising dataset | |
| 3. | 21/05/24 | Practical-2.1 Logistic Regression on Placement dataset | |
| 4. | 21/05/24 | Practical-2.2 Logistic Regression on iris dataset with Cross Validation and Kfold | |
| 5. | 21/05/24 | Practical-3.1_Lasso_Regression on Boston Housing dataset | |
| 6. | 21/05/24 | Practical-3.2_Linear_Regression_standard_scalar_california_housing | |
| 7. | 27/05/24 | Practical-3.3_ridge_Regression_standard_scalar_california_housing | |
| 8. | 27/05/24 | Practical-3.4_Lasso_Regression_standard_scalar_california_housing | |
| 9. | 28/05/24 | Practical-4_KNN_Regressor on Boston Housing dataset | |
| 10. | 28/05/24 | Practical-5_KNN_Classifier on iris dataset | |
| 11. | 30/05/24 | Practical-6_K-means on Student Results dataset | |
| 12. | 30/05/24 | Practical-7_SVM_classifier on diabetes dataset | |
| 13 | 06/06/24 | Practical-8_Ensemble_Bagging on Breast Cancer dataset | |

Practical-1.1_linear_regression_study_dataset

```
[ ]: !pip install scikit-learn pandas numpy
```

```
[1]: import pandas as pd
df = pd.read_csv("./dataset/Study.csv")
df
```

```
[1]:
```

| | Hours | Scores |
|----|-------|--------|
| 0 | 2.5 | 21 |
| 1 | 5.1 | 47 |
| 2 | 3.2 | 27 |
| 3 | 8.5 | 75 |
| 4 | 3.5 | 30 |
| 5 | 1.5 | 20 |
| 6 | 9.2 | 88 |
| 7 | 5.5 | 60 |
| 8 | 8.3 | 81 |
| 9 | 2.7 | 25 |
| 10 | 7.7 | 85 |
| 11 | 5.9 | 62 |
| 12 | 4.5 | 41 |
| 13 | 3.3 | 42 |
| 14 | 1.1 | 17 |
| 15 | 8.9 | 95 |
| 16 | 2.5 | 30 |
| 17 | 1.9 | 24 |
| 18 | 6.1 | 67 |
| 19 | 7.4 | 69 |
| 20 | 2.7 | 30 |
| 21 | 4.8 | 54 |
| 22 | 3.8 | 35 |
| 23 | 6.9 | 76 |
| 24 | 7.8 | 86 |

```
[2]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
#
```

```
---  -----  -----  ---
0  Hours    25 non-null    float64
1  Scores   25 non-null    int64
dtypes: float64(1), int64(1)
memory usage: 532.0 bytes
```

```
[3]: x = df.iloc[:, :-1]
     y = df.iloc[:, -1]
```

```
[4]: x
```

```
[4]:      Hours
0      2.5
1      5.1
2      3.2
3      8.5
4      3.5
5      1.5
6      9.2
7      5.5
8      8.3
9      2.7
10     7.7
11     5.9
12     4.5
13     3.3
14     1.1
15     8.9
16     2.5
17     1.9
18     6.1
19     7.4
20     2.7
21     4.8
22     3.8
23     6.9
24     7.8
```

```
[5]: y
```

```
[5]: 0      21
     1      47
     2      27
     3      75
     4      30
     5      20
     6      88
     7      60
```

```
8      81
9      25
10     85
11     62
12     41
13     42
14     17
15     95
16     30
17     24
18     67
19     69
20     30
21     54
22     35
23     76
24     86
Name: Scores, dtype: int64
```

```
[6]: from sklearn.model_selection import train_test_split
```

```
[7]: xtrain, xtest, ytrain, ytest = train_test_split(x,y,test_size=0.
      ↪2,random_state=1)
```

```
[8]: ytrain
```

```
[8]: 10      85
      18      67
      19      69
      4      30
      2      27
      20      30
      6      88
      7      60
      22      35
      1      47
      16      30
      0      21
      15      95
      24      86
      23      76
      9      25
      8      81
      12      41
      11      62
      5      20
Name: Scores, dtype: int64
```

```
[9]: from sklearn.linear_model import LinearRegression
```

```
[10]: lr = LinearRegression()
```

```
[11]: lr.fit(xtrain, ytrain)
```

```
[11]: LinearRegression()
```

```
[12]: predictions = lr.predict(xtest)
```

Beta 0

```
[14]: lr.intercept_
```

```
[14]: -1.5369573315500702
```

Beta 1

```
[15]: lr.coef_
```

```
[15]: array([10.46110829])
```

```
[16]: print(ytest)
      print(predictions)
```

```
14    17
13    42
17    24
3     75
21    54
Name: Scores, dtype: int64
[ 9.97026179 32.98470004 18.33914843 87.38246316 48.67636248]
```

```
[17]: from sklearn.metrics import mean_absolute_error, r2_score
```

```
[18]: mean_absolute_error(ytest, predictions)
```

```
[18]: 7.882398086270432
```

R2 Score should be close to 1

```
[19]: r2_score(ytest, predictions)
```

```
[19]: 0.8421031525243527
```

Practical-1.2_linear_regression_Advertising_dataset

```
[ ]: !pip install numpy pandas scikit-learn matplotlib
```

```
[2]: import pandas as pd
```

```
[3]: df = pd.read_csv('./dataset/Advertising.csv')
```

```
[4]: df
```




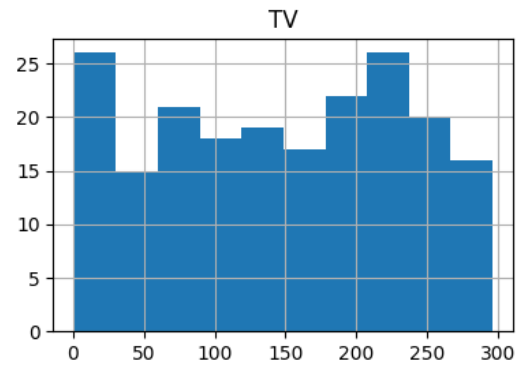

```
[4]:
```

| | ID | TV | Radio | Newspaper | Sales |
|-----|-----|-------|-------|-----------|-------|
| 0 | 1 | 230.1 | 37.8 | 69.2 | 22.1 |
| 1 | 2 | 44.5 | 39.3 | 45.1 | 10.4 |
| 2 | 3 | 17.2 | 45.9 | 69.3 | 9.3 |
| 3 | 4 | 151.5 | 41.3 | 58.5 | 18.5 |
| 4 | 5 | 180.8 | 10.8 | 58.4 | 12.9 |
| ... | ... | ... | ... | ... | ... |
| 195 | 196 | 38.2 | 3.7 | 13.8 | 7.6 |
| 196 | 197 | 94.2 | 4.9 | 8.1 | 9.7 |
| 197 | 198 | 177.0 | 9.3 | 6.4 | 12.8 |
| 198 | 199 | 283.6 | 42.0 | 66.2 | 25.5 |
| 199 | 200 | 232.1 | 8.6 | 8.7 | 13.4 |

[200 rows x 5 columns]

```
[5]: df.hist(figsize = (10,10))
```

```
[5]: array([[<Axes: title={'center': 'ID'}>, <Axes: title={'center': 'TV'}>],  
        [<Axes: title={'center': 'Radio'}>,  
         <Axes: title={'center': 'Newspaper'}>],  
        [<Axes: title={'center': 'Sales'}>, <Axes: >]], dtype=object)
```

```
[6]: x = df.iloc[:, :-1]
      x
```

```
[6]:
```

| | ID | TV | Radio | Newspaper |
|-----|-----|-------|-------|-----------|
| 0 | 1 | 230.1 | 37.8 | 69.2 |
| 1 | 2 | 44.5 | 39.3 | 45.1 |
| 2 | 3 | 17.2 | 45.9 | 69.3 |
| 3 | 4 | 151.5 | 41.3 | 58.5 |
| 4 | 5 | 180.8 | 10.8 | 58.4 |
| .. | ... | ... | ... | ... |
| 195 | 196 | 38.2 | 3.7 | 13.8 |
| 196 | 197 | 94.2 | 4.9 | 8.1 |

```

197  198  177.0    9.3    6.4
198  199  283.6   42.0   66.2
199  200  232.1    8.6    8.7

```

[200 rows x 4 columns]

```
[7]: y = df.iloc[:, -1]
      y
```

```
[7]: 0    22.1
      1    10.4
      2     9.3
      3    18.5
      4    12.9
```

...

```

195     7.6
196     9.7
197    12.8
198    25.5
199    13.4

```

Name: Sales, Length: 200, dtype: float64

```
[8]: from sklearn.model_selection import train_test_split
      xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.3,random_state=2)
```

```
[9]: from sklearn.linear_model import LinearRegression
      lr = LinearRegression()
```

```
[10]: lr.fit(xtrain,ytrain)
```

```
[10]: LinearRegression()
```

```
[11]: predict1 = lr.predict(xtest)
```

```
[12]: predict1
```

```
[12]: array([13.9532874 ,  9.66868515,  6.8184482 , 15.17477774, 18.17445759,
            15.71712804,  7.57102926, 20.49761622, 13.19168021, 17.33321225,
            11.14787534, 19.48405747,  9.21291503, 10.79916797, 13.87449153,
            12.82277648,  9.32688188, 18.03502362, 16.4741683 , 18.82432455,
            16.92036691, 16.04687481, 12.02316254, 12.14339122, 14.87243899,
            12.08358743, 15.49759505,  8.08839314, 16.76959963, 13.87908137,
            16.20151428, 17.08595853, 13.05758661, 13.02632284,  8.91762832,
            11.01315087, 21.97837505, 19.89971271, 15.98938391, 20.23559933,
            21.12563546, 17.17066585, 21.13093152, 15.04421578, 19.66106778,
            18.74322737, 17.58850441, 10.34247317,  9.60222257, 13.00297864,
            12.59823575, 14.39336676, 17.46356839, 16.99371718,  8.55587416,
            17.10336367,  9.13263193,  4.15865759,  7.67835429, 24.67819442])
```

```
[13]: from sklearn.metrics import r2_score, mean_absolute_error
```

```
[14]: print(r2_score(ytest, predict1))
```

0.8021354391516504

```
[15]: print(mean_absolute_error(ytest, predict1))
```

1.5264827355298254

Practical-2.1_Logistic_Regression_placement_dataset

```
[1]: import pandas as pd
```

```
[2]: df = pd.read_csv("./dataset/placement.csv")
```

```
[3]: df
```

```
[3]:
```

| | cgpa | placement_exam_marks | placed |
|-----|------|----------------------|--------|
| 0 | 7.19 | 26.0 | 1 |
| 1 | 7.46 | 38.0 | 1 |
| 2 | 7.54 | 40.0 | 1 |
| 3 | 6.42 | 8.0 | 1 |
| 4 | 7.23 | 17.0 | 0 |
| .. | ... | ... | ... |
| 995 | 8.87 | 44.0 | 1 |
| 996 | 9.12 | 65.0 | 1 |
| 997 | 4.89 | 34.0 | 0 |
| 998 | 8.62 | 46.0 | 1 |
| 999 | 4.90 | 10.0 | 1 |

[1000 rows x 3 columns]

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1000 entries, 0 to 999  
Data columns (total 3 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   cgpa                   1000 non-null   float64  
1   placement_exam_marks  1000 non-null   float64  
2   placed                 1000 non-null   int64  
dtypes: float64(2), int64(1)  
memory usage: 23.6 KB
```

```
[5]: df.shape
```

```
[5]: (1000, 3)
```

```
[6]: df.head()
```

```
[6]:      cgpa    placement_exam_marks    placed
0    7.19                26.0            1
1    7.46                38.0            1
2    7.54                40.0            1
3    6.42                 8.0            1
4    7.23                17.0            0
```

```
[7]: df['placed'].unique()
```

```
[7]: array([1, 0])
```

```
[8]: X = df.iloc[:, :-1]
```

```
[9]: y = df.iloc[:, -1]
```

```
10]: from sklearn.model_selection import train_test_split
```

```
11]: xtrain,xtest,ytrain,ytest = train_test_split(X,y,test_size=0.25, random_state=1)
```

```
12]: from sklearn.linear_model import LogisticRegression
```

```
13]: classifier = LogisticRegression()
```

```
14]: classifier.fit(xtrain, ytrain)
```

```
14]: LogisticRegression()
```

```
15]: predictions = classifier.predict(xtest)
```

```
16]: probability = classifier.predict_proba(xtest)
```

```
17]: probability[:10]
```

```
17]: array([[0.57188237, 0.42811763],
            [0.58207783, 0.41792217],
            [0.54254942, 0.45745058],
            [0.52320624, 0.47679376],
            [0.55570161, 0.44429839],
            [0.51994635, 0.48005365],
            [0.49926819, 0.50073181],
            [0.51124366, 0.48875634],
            [0.55763857, 0.44236143],
            [0.497788, 0.502212]])
```

```
18]: predictions
```

```
[18]: array([0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
           0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
```

```
0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
0, 0, 1, 0, 0, 0, 0, 0, 0])
```

```
[19]: from sklearn.metrics import accuracy_score
```

```
[20]: accuracy_score(ytest, predictions)
```

```
[20]: 0.48
```

Practical-2.2_Logistic_Regression_iris_cv_kfold

```
[1]: from sklearn.datasets import load_iris
```

```
[2]: from sklearn.model_selection import cross_val_score, KFold
```

```
[3]: from sklearn.linear_model import LogisticRegression
```

```
[4]: iris = load_iris()
```

```
[5]: x = iris.data
```

```
[14]: x[:10]
```

```
[14]: array([[5.1, 3.5, 1.4, 0.2],  
            [4.9, 3. , 1.4, 0.2],  
            [4.7, 3.2, 1.3, 0.2],  
            [4.6, 3.1, 1.5, 0.2],  
            [5. , 3.6, 1.4, 0.2],  
            [5.4, 3.9, 1.7, 0.4],  
            [4.6, 3.4, 1.4, 0.3],  
            [5. , 3.4, 1.5, 0.2],  
            [4.4, 2.9, 1.4, 0.2],  
            [4.9, 3.1, 1.5, 0.1]])
```

```
[7]: y = iris.target
```

```
[8]: y
```

```
[8]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
            0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
            1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
            2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])
```

```
[9]: lr = LogisticRegression()
```

```
[10]: k_fold = KFold(n_splits=7)
```

```
[11]: score = cross_val_score(lr, x, y, cv=k_fold)
```

```
[12]: print("Cross Validation Scores:{}".format(score))
```

```
Cross Validation Scores:[1.          1.          1.          0.80952381  0.95238095
0.85714286
0.9047619 ]
```

```
[13]: print("Average Cross Validation score:{}".format(score.mean()))
```

```
Average Cross Validation score:0.9319727891156463
```


Practical-3.1_Lasso_Regression_housing-dataset

```
[ ]: !pip install numpy pandas scikit-learn
```

```
[1]: import pandas as pd
df = pd.read_csv("./dataset/boston-housing-dataset.csv")

# First 3
df.head(3)

# Random 3
df.sample(3)

# Last 3
df.tail(3)
```

```
[1]:
```

| | CRIM | ZN | INDUS | CHAS | NX | RM | AGE | DIS | RAD | TAX | \ |
|-----|---------|-----|-------|------|-------|-------|------|--------|-----|-------|---|
| 503 | 0.06076 | 0.0 | 11.93 | 0 | 0.573 | 6.976 | 91.0 | 2.1675 | 1 | 273.0 | |
| 504 | 0.10959 | 0.0 | 11.93 | 0 | 0.573 | 6.794 | 89.3 | 2.3889 | 1 | 273.0 | |
| 505 | 0.04741 | 0.0 | 11.93 | 0 | 0.573 | 6.030 | 80.8 | 2.5050 | 1 | 273.0 | |

| | PTRATIO | B | LSTAT | MEDV |
|-----|---------|--------|-------|------|
| 503 | 21.0 | 396.90 | 5.64 | 23.9 |
| 504 | 21.0 | 393.45 | 6.48 | 22.0 |
| 505 | 21.0 | 396.90 | 7.88 | 11.9 |

```
[2]: x = df.iloc[:, :-1]
x.shape
```

```
[2]: (506, 13)
```

```
[3]: y = df.iloc[:, -1]
y
```

```
[3]: 0    24.0
1    21.6
2    34.7
3    33.4
4    36.2
...
```

```
501    22.4
502    20.6
503    23.9
504    22.0
505    11.9
Name: MEDV, Length: 506, dtype: float64
```

```
[4]: from sklearn.linear_model import Lasso
model = Lasso()
```

```
[5]: from sklearn.model_selection import train_test_split
```

```
[6]: xtrain, xtest, ytrain, ytest = train_test_split(x,y,test_size = 0.25,
↳random_state = 1)
model.fit(xtrain, ytrain)
```

```
[6]: Lasso()
```

```
[7]: from sklearn.model_selection import RepeatedKFold
cv = RepeatedKFold(n_splits = 10, n_repeats=3, random_state=1)
```

```
[8]: from sklearn.metrics import r2_score
ypred = model.predict(xtest)
r2_score(ytest, ypred)
```

```
[8]: 0.662198077052326
```

```
[9]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_sc = sc.fit_transform(x)
xtrain, xtest, ytrain, ytest = train_test_split(x_sc,y,test_size = 0.25,
↳random_state=1)
model1 = Lasso()
params = {'alpha':[0.00001,0.0001,0.001,0.01]}
```

```
[10]: from sklearn.model_selection import GridSearchCV
search = GridSearchCV(model1, params, cv=cv)
result = search.fit(x_sc,y)
result.best_params_
```

```
[10]: {'alpha': 0.01}
```

```
[11]: model2 = Lasso(alpha=0.01)
model2.fit(xtrain,ytrain)
```

```
[11]: Lasso(alpha=0.01)
```

```
[12]: ypred2 = model2.predict(xtest)  
      r2_score(ytest, ypred2)
```

```
[12]: 0.7787372388293925
```

Practical-

3.2_Linear_Regression_standard_scalar_california_housing

```
[1]: !pip install numpy pandas scikit-learn
```

```
[1]: import pandas as pd
from sklearn.datasets import fetch_california_housing
```

```
[2]: boston = fetch_california_housing()
boston_df = pd.DataFrame(boston.data, columns=boston.feature_names)
boston_df["Price_House"] = boston.target
boston_df
```

```
[2]:
```

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude \ |
|-------|--------|----------|----------|-----------|------------|----------|------------|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 |
| 3 | 5.6431 | 52.0 | 5.817352 | 1.073059 | 558.0 | 2.547945 | 37.85 |
| 4 | 3.8462 | 52.0 | 6.281853 | 1.081081 | 565.0 | 2.181467 | 37.85 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 20635 | 1.5603 | 25.0 | 5.045455 | 1.133333 | 845.0 | 2.560606 | 39.48 |
| 20636 | 2.5568 | 18.0 | 6.114035 | 1.315789 | 356.0 | 3.122807 | 39.49 |
| 20637 | 1.7000 | 17.0 | 5.205543 | 1.120092 | 1007.0 | 2.325635 | 39.43 |
| 20638 | 1.8672 | 18.0 | 5.329513 | 1.171920 | 741.0 | 2.123209 | 39.43 |
| 20639 | 2.3886 | 16.0 | 5.254717 | 1.162264 | 1387.0 | 2.616981 | 39.37 |

| | Longitude | Price_House |
|-------|-----------|-------------|
| 0 | -122.23 | 4.526 |
| 1 | -122.22 | 3.585 |
| 2 | -122.24 | 3.521 |
| 3 | -122.25 | 3.413 |
| 4 | -122.25 | 3.422 |
| ... | ... | ... |
| 20635 | -121.09 | 0.781 |
| 20636 | -121.21 | 0.771 |
| 20637 | -121.22 | 0.923 |
| 20638 | -121.32 | 0.847 |
| 20639 | -121.24 | 0.894 |

[20640 rows x 9 columns]

```
[3]: x = boston_df.iloc[:, :-1]
```

```
[4]: x
```

```
[4]:
```

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | \ |
|-------|--------|----------|----------|-----------|------------|----------|----------|---|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | |
| 3 | 5.6431 | 52.0 | 5.817352 | 1.073059 | 558.0 | 2.547945 | 37.85 | |
| 4 | 3.8462 | 52.0 | 6.281853 | 1.081081 | 565.0 | 2.181467 | 37.85 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 20635 | 1.5603 | 25.0 | 5.045455 | 1.133333 | 845.0 | 2.560606 | 39.48 | |
| 20636 | 2.5568 | 18.0 | 6.114035 | 1.315789 | 356.0 | 3.122807 | 39.49 | |
| 20637 | 1.7000 | 17.0 | 5.205543 | 1.120092 | 1007.0 | 2.325635 | 39.43 | |
| 20638 | 1.8672 | 18.0 | 5.329513 | 1.171920 | 741.0 | 2.123209 | 39.43 | |
| 20639 | 2.3886 | 16.0 | 5.254717 | 1.162264 | 1387.0 | 2.616981 | 39.37 | |

| | Longitude |
|-------|-----------|
| 0 | -122.23 |
| 1 | -122.22 |
| 2 | -122.24 |
| 3 | -122.25 |
| 4 | -122.25 |
| ... | ... |
| 20635 | -121.09 |
| 20636 | -121.21 |
| 20637 | -121.22 |
| 20638 | -121.32 |
| 20639 | -121.24 |

[20640 rows x 8 columns]

```
[5]: y = boston_df.iloc[:, -1]
```

```
[6]: y
```

```
[6]:
```

| | |
|-------|-------|
| 0 | 4.526 |
| 1 | 3.585 |
| 2 | 3.521 |
| 3 | 3.413 |
| 4 | 3.422 |
| ... | ... |
| 20635 | 0.781 |
| 20636 | 0.771 |
| 20637 | 0.923 |

```
20638    0.847
20639    0.894
Name: Price_House, Length: 20640, dtype: float64
```

```
[7]: from sklearn.preprocessing import StandardScaler
     sc = StandardScaler()
     x_sc= sc.fit_transform(x)
```

```
[8]: x_sc
```

```
[8]: array([[ 2.34476576,  0.98214266,  0.62855945, ..., -0.04959654,
            1.05254828, -1.32783522],
           [ 2.33223796, -0.60701891,  0.32704136, ..., -0.09251223,
            1.04318455, -1.32284391],
           [ 1.7826994 ,  1.85618152,  1.15562047, ..., -0.02584253,
            1.03850269, -1.33282653],
           ...,
           [-1.14259331, -0.92485123, -0.09031802, ..., -0.0717345 ,
            1.77823747, -0.8237132 ],
           [-1.05458292, -0.84539315, -0.04021111, ..., -0.09122515,
            1.77823747, -0.87362627],
           [-0.78012947, -1.00430931, -0.07044252, ..., -0.04368215,
            1.75014627, -0.83369581]])
```

```
[9]: from sklearn.model_selection import train_test_split
     xtrain,xtest,ytrain,ytest = train_test_split(x_sc,y,test_size=0.
     ↪3,random_state=2)
```

```
[10]: from sklearn.linear_model import LinearRegression
     lr = LinearRegression()
```

```
[11]: lr.fit(xtrain,ytrain)
```

```
[11]: LinearRegression()
```

```
[12]: predict1 = lr.predict(xtest)
```

```
[13]: from sklearn.metrics import r2_score,mean_absolute_error
```

```
[14]: print(r2_score(ytest,predict1))
```

```
0.6015507891610434
```

```
[15]: print(mean_absolute_error(ytest,predict1))
```

```
0.5362588391493065
```

Practical-

3.3_ridge_Regression_standard_scalar_california_housing

```
[ ]: !pip install numpy pandas scikit-learn
```

```
[1]: import pandas as pd
from sklearn.datasets import fetch_california_housing
ds = fetch_california_housing()
df = pd.DataFrame(ds.data, columns=ds.feature_names)
df.head(3)
```

```
[1]:
```

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | \ |
|---|--------|----------|----------|-----------|------------|----------|----------|---|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | |

| | Longitude |
|---|-----------|
| 0 | -122.23 |
| 1 | -122.22 |
| 2 | -122.24 |

```
[2]: df["HousePrice"]=ds.target
df.head(3)
```

```
[2]:
```

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | \ |
|---|--------|----------|----------|-----------|------------|----------|----------|---|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | |

| | Longitude | HousePrice |
|---|-----------|------------|
| 0 | -122.23 | 4.526 |
| 1 | -122.22 | 3.585 |
| 2 | -122.24 | 3.521 |

```
[3]: X = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

```
[4]: X.head(3)
```

```
[4]: MedInc HouseAge AveRooms AveBedrms Population AveOccup Latitude \
0 8.3252 41.0 6.984127 1.023810 322.0 2.555556 37.88
1 8.3014 21.0 6.238137 0.971880 2401.0 2.109842 37.86
2 7.2574 52.0 8.288136 1.073446 496.0 2.802260 37.85

Longitude
0 -122.23
1 -122.22
2 -122.24
```

```
[5]: y.head(3)
```

```
[5]: 0 4.526
1 3.585
2 3.521
Name: HousePrice, dtype: float64
```

```
[6]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_sc = sc.fit_transform(X)
x_sc
```

```
[6]: array([[ 2.34476576,  0.98214266,  0.62855945, ..., -0.04959654,
           1.05254828, -1.32783522],
          [ 2.33223796, -0.60701891,  0.32704136, ..., -0.09251223,
           1.04318455, -1.32284391],
          [ 1.7826994 ,  1.85618152,  1.15562047, ..., -0.02584253,
           1.03850269, -1.33282653],
          ...,
          [-1.14259331, -0.92485123, -0.09031802, ..., -0.0717345 ,
           1.77823747, -0.8237132 ],
          [-1.05458292, -0.84539315, -0.04021111, ..., -0.09122515,
           1.77823747, -0.87362627],
          [-0.78012947, -1.00430931, -0.07044252, ..., -0.04368215,
           1.75014627, -0.83369581]])
```

```
[7]: from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x_sc, y, test_size = 0.3,
↪random_state=2)
```

```
[8]: from sklearn.linear_model import Ridge
rr = Ridge(alpha = 10)
```

```
[9]: rr.fit(xtrain,ytrain)
```

```
[9]: Ridge(alpha=10)
```

```
[10]: predict = rr.predict(xtest)
```



```
[11]: from sklearn.metrics import r2_score, mean_absolute_error  
      r2_score(ytest, predict)
```

```
[11]: 0.6014258698314037
```

```
[12]: mean_absolute_error(ytest, predict)
```

```
[12]: np.float64(0.536239634635214)
```

Practical-

3.4_Lasso_Regression_standard_scalar_california_housing

```
[ ]: !pip install numpy pandas scikit-learn
```

```
[1]: import pandas as pd
      from sklearn.datasets import fetch_california_housing
```

```
[2]: ds = fetch_california_housing()
      df = pd.DataFrame(ds.data, columns = ds.feature_names)
      df["HousePrice"] = ds.target

      df.head(3)
```

```
[2]:
```

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | \ |
|---|--------|----------|----------|-----------|------------|----------|----------|---|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | |

| | Longitude | HousePrice |
|---|-----------|------------|
| 0 | -122.23 | 4.526 |
| 1 | -122.22 | 3.585 |
| 2 | -122.24 | 3.521 |

```
[3]: X = df.iloc[:, :-1]
      y = df.iloc[:, -1]
```

```
[4]: X.head(3)
```

```
[4]:
```

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | \ |
|---|--------|----------|----------|-----------|------------|----------|----------|---|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | |

| | Longitude |
|---|-----------|
| 0 | -122.23 |
| 1 | -122.22 |
| 2 | -122.24 |

```
[5]: y.head(3)
```

```
[5]: 0    4.526
      1    3.585
      2    3.521
      Name: HousePrice, dtype: float64
```

```
[6]: from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
      x_sc = sc.fit_transform(X)

      x_sc
```

```
[6]: array([[ 2.34476576,  0.98214266,  0.62855945, ..., -0.04959654,
            1.05254828, -1.32783522],
      [ 2.33223796, -0.60701891,  0.32704136, ..., -0.09251223,
            1.04318455, -1.32284391],
      [ 1.7826994 ,  1.85618152,  1.15562047, ..., -0.02584253,
            1.03850269, -1.33282653],
      ...,
      [-1.14259331, -0.92485123, -0.09031802, ..., -0.0717345 ,
            1.77823747, -0.8237132 ],
      [-1.05458292, -0.84539315, -0.04021111, ..., -0.09122515,
            1.77823747, -0.87362627],
      [-0.78012947, -1.00430931, -0.07044252, ..., -0.04368215,
            1.75014627, -0.83369581]])
```

```
[7]: from sklearn.model_selection import train_test_split
      xtrain, xtest, ytrain, ytest = train_test_split(x_sc, y, test_size=0.3,
      ↪random_state=2)
```

```
[8]: from sklearn.linear_model import Lasso
      lr = Lasso(alpha=10)
```

```
[9]: lr.fit(xtrain, ytrain)
```

```
[9]: Lasso(alpha=10)
```

```
[10]: predict = lr.predict(xtest)
```

```
[11]: from sklearn.metrics import r2_score, mean_absolute_error
      r2_score(ytest, predict)
```

```
[11]: -0.00033321189864432554
```

```
[12]: mean_absolute_error(ytest, predict)
```

```
[12]: np.float64(0.9147942692371251)
```

Practical-4_KNN_Regressor-assignment-done

```
[1]: import pandas as pd
data = pd.read_csv("./dataset/BostonHousing.csv")
data.head(3)
```

```
[1]:      crim    zn  indus  chas    nox    rm   age    dis  rad  tax  ptratio  \
0  0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900   1  296    15.3
1  0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671   2  242    17.8
2  0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671   2  242    17.8

      b  lstat  medv
0  396.90   4.98  24.0
1  396.90   9.14  21.6
2  392.83   4.03  34.7
```

```
[2]: from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler(feature_range=(0,1))
```

```
[3]: x_og = data.iloc[:, :-1]
```

```
[4]: x = sc.fit_transform(x_og)
x
```

```
[4]: array([[0.00000000e+00, 1.80000000e-01, 6.78152493e-02, ...,
          2.87234043e-01, 1.00000000e+00, 8.96799117e-02],
          [2.35922539e-04, 0.00000000e+00, 2.42302053e-01, ...,
          5.53191489e-01, 1.00000000e+00, 2.04470199e-01],
          [2.35697744e-04, 0.00000000e+00, 2.42302053e-01, ...,
          5.53191489e-01, 9.89737254e-01, 6.34657837e-02],
          ...,
          [6.11892474e-04, 0.00000000e+00, 4.20454545e-01, ...,
          8.93617021e-01, 1.00000000e+00, 1.07891832e-01],
          [1.16072990e-03, 0.00000000e+00, 4.20454545e-01, ...,
          8.93617021e-01, 9.91300620e-01, 1.31070640e-01],
          [4.61841693e-04, 0.00000000e+00, 4.20454545e-01, ...,
          8.93617021e-01, 1.00000000e+00, 1.69701987e-01]])
```

```
[5]: x.shape
```

[5]: (506, 13)

```
[6]: y = data.iloc[:,-1]
      y.shape
```

[6]: (506,)

```
[7]: from sklearn.model_selection import train_test_split
      xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.25,
      ↪random_state = 1)
```

```
[8]: from sklearn.neighbors import KNeighborsRegressor
      model = KNeighborsRegressor(n_neighbors=3)
      model.fit(xtrain,ytrain)
```

[8]: KNeighborsRegressor(n_neighbors=3)

```
[9]: predictions = model.predict(xtest)
```

```
[10]: from sklearn.metrics import r2_score
      r2_score(ytest, predictions)
```

[10]: 0.8069614057252134

0.1 Testing with StandardScaler

```
[11]: from sklearn.preprocessing import StandardScaler
      ss = StandardScaler()
```

```
[12]: ss.fit(x_og)
      x = ss.transform(x_og)
```

```
[13]: xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.25,
      ↪random_state = 1)
```

0.2 Find best number of neighbors using gridSearchCV

```
[14]: from sklearn.model_selection import GridSearchCV
```

```
[15]: model2 = KNeighborsRegressor()
```

```
[16]: param = {"n_neighbors": [1,3,5,7,9,11],
              "weights": ["uniform", "distance"]}

      scoring = "neg_mean_squared_error"
```

```
[17]: search = GridSearchCV(model2, param, scoring = scoring, cv = 5)
      search.fit(xtrain, ytrain)
```

```
print(f"Best Params: {search.best_params_}")
```

Best Params: {'n_neighbors': 3, 'weights': 'distance'}

```
[18]: final_model = search.best_estimator_
```

```
[19]: predictions2 = final_model.predict(xtest)
```

```
from sklearn.metrics import r2_score  
r2_score(ytest, predictions2)
```

```
[19]: 0.8638489106891928
```

Practical-5_KNN_Classifier_assignment_done

```
[1]: from sklearn.datasets import load_iris
df = load_iris()
x = df.data
y = df.target
```

```
[2]: x
```

```
[2]: array([[5.1, 3.5, 1.4, 0.2],
 [4.9, 3. , 1.4, 0.2],
 [4.7, 3.2, 1.3, 0.2],
 [4.6, 3.1, 1.5, 0.2],
 [5. , 3.6, 1.4, 0.2],
 [5.4, 3.9, 1.7, 0.4],
 [4.6, 3.4, 1.4, 0.3],
 [5. , 3.4, 1.5, 0.2],
 [4.4, 2.9, 1.4, 0.2],
 [4.9, 3.1, 1.5, 0.1],
 [5.4, 3.7, 1.5, 0.2],
 [4.8, 3.4, 1.6, 0.2],
 [4.8, 3. , 1.4, 0.1],
 [4.3, 3. , 1.1, 0.1],
 [5.8, 4. , 1.2, 0.2],
 [5.7, 4.4, 1.5, 0.4],
 [5.4, 3.9, 1.3, 0.4],
 [5.1, 3.5, 1.4, 0.3],
 [5.7, 3.8, 1.7, 0.3],
 [5.1, 3.8, 1.5, 0.3],
 [5.4, 3.4, 1.7, 0.2],
 [5.1, 3.7, 1.5, 0.4],
 [4.6, 3.6, 1. , 0.2],
 [5.1, 3.3, 1.7, 0.5],
 [4.8, 3.4, 1.9, 0.2],
 [5. , 3. , 1.6, 0.2],
 [5. , 3.4, 1.6, 0.4],
 [5.2, 3.5, 1.5, 0.2],
 [5.2, 3.4, 1.4, 0.2],
 [4.7, 3.2, 1.6, 0.2],
 [4.8, 3.1, 1.6, 0.2],
```

[5.4, 3.4, 1.5, 0.4],
 [5.2, 4.1, 1.5, 0.1],
 [5.5, 4.2, 1.4, 0.2],
 [4.9, 3.1, 1.5, 0.2],
 [5. , 3.2, 1.2, 0.2],
 [5.5, 3.5, 1.3, 0.2],
 [4.9, 3.6, 1.4, 0.1],
 [4.4, 3. , 1.3, 0.2],
 [5.1, 3.4, 1.5, 0.2],
 [5. , 3.5, 1.3, 0.3],
 [4.5, 2.3, 1.3, 0.3],
 [4.4, 3.2, 1.3, 0.2],
 [5. , 3.5, 1.6, 0.6],
 [5.1, 3.8, 1.9, 0.4],
 [4.8, 3. , 1.4, 0.3],
 [5.1, 3.8, 1.6, 0.2],
 [4.6, 3.2, 1.4, 0.2],
 [5.3, 3.7, 1.5, 0.2],
 [5. , 3.3, 1.4, 0.2],
 [7. , 3.2, 4.7, 1.4],
 [6.4, 3.2, 4.5, 1.5],
 [6.9, 3.1, 4.9, 1.5],
 [5.5, 2.3, 4. , 1.3],
 [6.5, 2.8, 4.6, 1.5],
 [5.7, 2.8, 4.5, 1.3],
 [6.3, 3.3, 4.7, 1.6],
 [4.9, 2.4, 3.3, 1.],
 [6.6, 2.9, 4.6, 1.3],
 [5.2, 2.7, 3.9, 1.4],
 [5. , 2. , 3.5, 1.],
 [5.9, 3. , 4.2, 1.5],
 [6. , 2.2, 4. , 1.],
 [6.1, 2.9, 4.7, 1.4],
 [5.6, 2.9, 3.6, 1.3],
 [6.7, 3.1, 4.4, 1.4],
 [5.6, 3. , 4.5, 1.5],
 [5.8, 2.7, 4.1, 1.],
 [6.2, 2.2, 4.5, 1.5],
 [5.6, 2.5, 3.9, 1.1],
 [5.9, 3.2, 4.8, 1.8],
 [6.1, 2.8, 4. , 1.3],
 [6.3, 2.5, 4.9, 1.5],
 [6.1, 2.8, 4.7, 1.2],
 [6.4, 2.9, 4.3, 1.3],
 [6.6, 3. , 4.4, 1.4],
 [6.8, 2.8, 4.8, 1.4],
 [6.7, 3. , 5. , 1.7],

[6. , 2.9, 4.5, 1.5],
[5.7, 2.6, 3.5, 1.],
[5.5, 2.4, 3.8, 1.1],
[5.5, 2.4, 3.7, 1.],
[5.8, 2.7, 3.9, 1.2],
[6. , 2.7, 5.1, 1.6],
[5.4, 3. , 4.5, 1.5],
[6. , 3.4, 4.5, 1.6],
[6.7, 3.1, 4.7, 1.5],
[6.3, 2.3, 4.4, 1.3],
[5.6, 3. , 4.1, 1.3],
[5.5, 2.5, 4. , 1.3],
[5.5, 2.6, 4.4, 1.2],
[6.1, 3. , 4.6, 1.4],
[5.8, 2.6, 4. , 1.2],
[5. , 2.3, 3.3, 1.],
[5.6, 2.7, 4.2, 1.3],
[5.7, 3. , 4.2, 1.2],
[5.7, 2.9, 4.2, 1.3],
[6.2, 2.9, 4.3, 1.3],
[5.1, 2.5, 3. , 1.1],
[5.7, 2.8, 4.1, 1.3],
[6.3, 3.3, 6. , 2.5],
[5.8, 2.7, 5.1, 1.9],
[7.1, 3. , 5.9, 2.1],
[6.3, 2.9, 5.6, 1.8],
[6.5, 3. , 5.8, 2.2],
[7.6, 3. , 6.6, 2.1],
[4.9, 2.5, 4.5, 1.7],
[7.3, 2.9, 6.3, 1.8],
[6.7, 2.5, 5.8, 1.8],
[7.2, 3.6, 6.1, 2.5],
[6.5, 3.2, 5.1, 2.],
[6.4, 2.7, 5.3, 1.9],
[6.8, 3. , 5.5, 2.1],
[5.7, 2.5, 5. , 2.],
[5.8, 2.8, 5.1, 2.4],
[6.4, 3.2, 5.3, 2.3],
[6.5, 3. , 5.5, 1.8],
[7.7, 3.8, 6.7, 2.2],
[7.7, 2.6, 6.9, 2.3],
[6. , 2.2, 5. , 1.5],
[6.9, 3.2, 5.7, 2.3],
[5.6, 2.8, 4.9, 2.],
[7.7, 2.8, 6.7, 2.],
[6.3, 2.7, 4.9, 1.8],
[6.7, 3.3, 5.7, 2.1],

```
[7.2, 3.2, 6. , 1.8],
[6.2, 2.8, 4.8, 1.8],
[6.1, 3. , 4.9, 1.8],
[6.4, 2.8, 5.6, 2.1],
[7.2, 3. , 5.8, 1.6],
[7.4, 2.8, 6.1, 1.9],
[7.9, 3.8, 6.4, 2. ],
[6.4, 2.8, 5.6, 2.2],
[6.3, 2.8, 5.1, 1.5],
[6.1, 2.6, 5.6, 1.4],
[7.7, 3. , 6.1, 2.3],
[6.3, 3.4, 5.6, 2.4],
[6.4, 3.1, 5.5, 1.8],
[6. , 3. , 4.8, 1.8],
[6.9, 3.1, 5.4, 2.1],
[6.7, 3.1, 5.6, 2.4],
[6.9, 3.1, 5.1, 2.3],
[5.8, 2.7, 5.1, 1.9],
[6.8, 3.2, 5.9, 2.3],
[6.7, 3.3, 5.7, 2.5],
[6.7, 3. , 5.2, 2.3],
[6.3, 2.5, 5. , 1.9],
[6.5, 3. , 5.2, 2. ],
[6.2, 3.4, 5.4, 2.3],
[5.9, 3. , 5.1, 1.8]])
```

```
[3]: y
```

```
[3]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])
```

```
[4]: from sklearn.model_selection import train_test_split
      xtrain, xtest, ytrain, ytest = train_test_split(x,y,test_size=0.2, random_state=
      ↪ 1)
```

```
[5]: from sklearn.neighbors import KNeighborsClassifier
      model = KNeighborsClassifier()
      model.fit(xtrain,ytrain)
```

```
[5]: KNeighborsClassifier()
```

```
[6]: predictions = model.predict(xtest)
```

```
[7]: from sklearn.metrics import accuracy_score
accuracy_score(ytest, predictions)
```

[7]: 1.0

0.1 Find best number of neighbors using gridSearchCV

```
[8]: from sklearn.model_selection import GridSearchCV
```

```
[9]: model2 = KNeighborsClassifier()
```

```
[10]: param = {"n_neighbors" : [1,3,5],
              "weights" : ["uniform", "distance"]}
scoring = "neg_mean_squared_error"
```

```
[11]: search = GridSearchCV(model2, param, scoring = scoring, cv = 5)
search.fit(xtrain, ytrain)
print(f"Best Params: {search.best_params_}")
```

Best Params: {'n_neighbors': 3, 'weights': 'uniform'}

```
[12]: final_model = search.best_estimator_
```

```
[13]: predictions2 = final_model.predict(xtest)
```

```
[14]: from sklearn.metrics import r2_score
r2_score(ytest, predictions2)
```

[14]: 1.0

Practical-6_K-means-R

```
[1]: library(ggplot2)
      library(cluster)
```

```
[3]: df = as.data.frame(read.csv("./dataset/marks.csv"))
      df
```


A data.frame: 17 × 4

| | Roll_no <int> | English <int> | Maths <int> | Science <int> |
|----|------------------|------------------|----------------|------------------|
| 1 | 99 | 100 | 100 | |
| 2 | 98 | 99 | 97 | |
| 3 | 92 | 9 | 96 | |
| 4 | 95 | 92 | 94 | |
| 5 | 90 | 100 | 96 | |
| 6 | 80 | 75 | 82 | |
| 7 | 75 | 83 | 80 | |
| 8 | 72 | 73 | 74 | |
| 9 | 71 | 82 | 76 | |
| 10 | 73 | 74 | 76 | |
| 11 | 34 | 32 | 28 | |
| 12 | 26 | 28 | 30 | |
| 13 | 32 | 30 | 31 | |
| 14 | 98 | 97 | 98 | |
| 15 | 30 | 29 | 29 | |
| 16 | 78 | 75 | 78 | |
| 17 | 100 | 99 | 100 | |

```
[4]: kmdata = df[,2:4]
      kmdata
```

| | English <int> | Maths <int> | Science <int> |
|----------------------|------------------|----------------|------------------|
| | 99 | 100 | 100 |
| | 98 | 99 | 97 |
| | 92 | 9 | 96 |
| | 95 | 92 | 94 |
| | 90 | 100 | 96 |
| | 80 | 75 | 82 |
| | 75 | 83 | 80 |
| A data.frame: 17 × 3 | 72 | 73 | 74 |
| | 71 | 82 | 76 |
| | 73 | 74 | 76 |
| | 34 | 32 | 28 |
| | 26 | 28 | 30 |
| | 32 | 30 | 31 |
| | 98 | 97 | 98 |
| | 30 | 29 | 29 |
| | 78 | 75 | 78 |
| | 100 | 99 | 100 |

```
[5]: wss = numeric(8)
for (k in 1:8) wss[k] = sum(kmeans(kmdata,centers = k, nstart = 25)$withinss)
plot(1:8,wss,type="b",xlab="no of clusters", ylab = "wssvalues")
```



```
[6]: final_model = kmeans(kmdata,3,nstart = 25)
final_model
```

K-means clustering with 3 clusters of sizes 4, 1, 12

Cluster means:

| | English | Maths | Science |
|---|---------|----------|----------|
| 1 | 30.50 | 29.75000 | 29.50000 |
| 2 | 92.00 | 9.00000 | 96.00000 |
| 3 | 85.75 | 87.41667 | 87.58333 |

Clustering vector:

```
[1] 3 3 2 3 3 3 3 3 3 3 1 1 1 3 1 3 3
```

Within cluster sum of squares by cluster:

```
[1] 48.750 0.000 4254.083  
(between_SS / total_SS = 88.8 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

Practical-7_SVM_classifier

```
[1]: import pandas as pd
```

```
[4]: df = pd.read_csv("./dataset/diabetes.csv")
```

```
[5]: df.head(3)
```

```
[5]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|---|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |

```
[6]: df.sample(2)
```

```
[6]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|-----|-------------|---------|---------------|---------------|---------|------|---|
| 442 | 4 | 117 | 64 | 27 | 120 | 33.2 | |
| 311 | 0 | 106 | 70 | 37 | 148 | 39.4 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|-----|--------------------------|-----|---------|
| 442 | 0.230 | 24 | 0 |
| 311 | 0.605 | 22 | 0 |

```
[7]: df.shape
```

```
[7]: (768, 9)
```

```
[8]: df["Outcome"].value_counts()
```

```
[8]: Outcome
0    500
1    268
Name: count, dtype: int64
```



```
[9]: X = df.iloc[:, :-1]
X.shape
```

```
[9]: (768, 8)
```

```
[10]: y = df.iloc[:, -1]
y.shape
```

```
[10]: (768,)
```

```
[11]: X
```

```
[11]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0                6      148             72             35         0  33.6
1                1       85             66             29         0  26.6
2                8      183             64              0         0  23.3
3                1       89             66             23        94  28.1
4                0      137             40             35       168  43.1
..            ...    ...             ...             ...    ...    ...
763             10      101             76             48       180  32.9
764              2      122             70             27         0  36.8
765              5      121             72             23       112  26.2
766              1      126             60              0         0  30.1
767              1       93             70             31         0  30.4
```

```
      DiabetesPedigreeFunction  Age
0                0.627      50
1                0.351      31
2                0.672      32
3                0.167      21
4                2.288      33
..            ...    ...
763             0.171      63
764             0.340      27
765             0.245      30
766             0.349      47
767             0.315      23
```

```
[768 rows x 8 columns]
```

```
[12]: y
```

```
[12]: 0      1
1      0
2      1
3      0
4      1
..
```

```
763    0
764    0
765    0
766    1
767    0
```

Name: Outcome, Length: 768, dtype: int64

```
[13]: from sklearn.model_selection import train_test_split
      xtrain, xtest, ytrain, ytest = train_test_split(X, y, test_size = 0.25,
      ↪random_state=1)
```

```
[14]: xtrain.shape, xtest.shape, ytrain.shape, ytest.shape
```

```
[14]: ((576, 8), (192, 8), (576,), (192,))
```

```
[15]: from sklearn.svm import SVC

      model1 = SVC(kernel="linear", degree=2, coef0=1.5)
      # default will be linear kernel

      model1.fit(xtrain, ytrain)
```

```
[15]: SVC(coef0=1.5, degree=2, kernel='linear')
```

```
[16]: predictions = model1.predict(xtest)
```

```
[17]: from sklearn.metrics import accuracy_score, r2_score, confusion_matrix

      accuracy_score(ytest, predictions)
```

```
[17]: 0.78125
```

```
[18]: confusion_matrix(ytest, predictions)
```

```
[18]: array([[108, 15],
      [ 27, 42]])
```

```
[19]: model2 = SVC(kernel="poly", degree=2, coef0=1.5)
      model2.fit(xtrain, ytrain)
```

```
[19]: SVC(coef0=1.5, degree=2, kernel='poly')
```

```
[20]: predictions2 = model2.predict(xtest)
      accuracy_score(ytest, predictions2)
```

```
[20]: 0.796875
```

```
[21]: confusion_matrix(ytest, predictions2)
```

```
[21]: array([[113, 10],  
           [ 29, 40]])
```

```
[22]: model3 = SVC(kernel="rbf", degree=2, coef0=1.5)  
      model3.fit(xtrain, ytrain)
```

```
[22]: SVC(coef0=1.5, degree=2)
```

```
[23]: predictions3 = model3.predict(xtest)  
      accuracy_score(ytest, predictions3)
```

```
[23]: 0.7708333333333334
```

Practical-8_Ensemble_Bagging

```
[1]: import numpy as np
      from sklearn.datasets import load_breast_cancer
      from sklearn.model_selection import train_test_split
```

```
[2]: from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import LogisticRegression
      from sklearn.pipeline import make_pipeline
      from sklearn.ensemble import BaggingClassifier
```

```
[3]: df = load_breast_cancer()
      x=df.data
      y=df.target
```

```
[4]: xtrain, xtest, ytrain, ytest = train_test_split(x,y,test_size=0.
      ↪25,random_state=1)
```

```
[5]: pipeline= make_pipeline(StandardScaler(),LogisticRegression(random_state=1))
```

```
[6]: pipeline.fit(xtrain,ytrain)
```

```
[6]: Pipeline(steps=[('standardscaler', StandardScaler()),
                      ('logisticregression', LogisticRegression(random_state=1))])
```

```
[7]: print("Score:",pipeline.score(xtest,ytest))
      pred1=pipeline.predict(xtest)
      from sklearn.metrics import accuracy_score
      print("Accuracy Score:",accuracy_score(ytest,pred1))
```

Score: 0.9790209790209791

Accuracy Score: 0.9790209790209791

```
[8]: bgclassifier=
      ↪BaggingClassifier(estimator=pipeline,n_estimators=100,max_features=10,max_samples=100,
      ↪random_state=1)
```

```
[9]: bgclassifier.fit(xtrain,ytrain)
```

```
[9]: BaggingClassifier(estimator=Pipeline(steps=[('standardscaler',
                                                StandardScaler()),
```

```
                                ('logisticregression',  
    LogisticRegression(random_state=1))]),  
                                max_features=10, max_samples=100, n_estimators=100,  
                                random_state=1)
```

```
[10]: print(bgclassifier.score(xtest,ytest))
```

```
0.958041958041958
```

```
[ ]:
```